

Apache Flink for Stream Processing (20 hrs)

By Dr. Vishwanath Rao

Module 1: Introduction to Apache Flink

- Overview of Stream Processing
 - Differences between batch and stream processing
 - Use cases and applications of Flink
- Introduction to Apache Flink
 - Architecture and key components
 - Flink ecosystem

Module 2: Flink Streaming Fundamentals

- Introduction to Flink Streaming
 - DataStream API vs. DataSet API
 - Creating and processing DataStreams
- Data Sources and Sinks
 - Reading data from sources (Kafka, files, etc.)
 - Writing data to sinks (databases, file systems)

Module 3: Windowing and Process Functions

- Understanding Windowing Concepts
 - Types of windows: Tumbling, Sliding, Session
 - Use cases for windowing
- Implementing Aggregations
 - Common aggregation functions
 - Custom aggregators
- Process Functions
 - Keyed and non-keyed process functions
 - Use cases and examples

Module 4: Watermarking and Handling Late-Coming Events

- Introduction to Watermarking
 - What are watermarks and their purpose?
 - Types of watermarks (periodic, punctuated)
- Handling Late Events
 - Strategies for dealing with late arrivals
 - Use cases for late-event handling

Module 5: State Management and Optimization

- Understanding State in Flink

- Types of state: Keyed state vs. Operator state
- Use cases for stateful processing
- State Optimization Techniques
 - Best practices for state management
 - Memory management and tuning

Module 6: Observability and Monitoring

- Importance of Observability in Stream Processing
 - Key metrics to monitor in Flink
- Tools and Techniques for Monitoring
 - Using the Flink Web UI
 - Integrating with external monitoring tools (Prometheus, Grafana)

Module 7: Checkpointing and Savepoints

- Understanding Checkpointing
 - What are checkpoints and their purpose?
 - Configuring checkpoints for fault tolerance
- Savepoints
 - Difference between checkpoints and savepoints
 - How to create and manage savepoints

Module 8: Connectors in Flink

- Overview of Flink Connectors
 - Common connectors: Kafka, JDBC, Files, etc.
- Implementing Custom Connectors
 - Creating and configuring custom connectors
 - Best practices for connector development

Module 9: Parallelism in Flink

- Understanding Parallelism
 - What is parallelism and why is it important?
 - Configuring parallelism for tasks and operators
- Strategies for Optimizing Parallelism
 - Load balancing and resource allocation
 - Case studies of parallel processing

Module 10: Debugging and Metrics

- Debugging Flink Applications
 - Common debugging techniques and tools
 - Analyzing errors and exceptions

- Collecting and Analyzing Metrics
 - Instrumentation of Flink applications
 - Using metrics for performance tuning